End-to-end latency to start a pod
(cold start='.' - warm start='0')

Legend:
- pod-scheduling
- make-pod-sandbox
- image-pull
- create-container
- start-container
- execute-job

Y-axis: Time [s]
X-axis: Baseline (coco, docker, coco-nydus, kata)